



VISUALIZING NETWORK DATA

The cliché goes that the world is an increasingly interconnected place, and the connections between different entities are often best represented with a graph. Graphs are comprised of vertices (also often called "nodes") and edges connecting those nodes. In this assignment, we will learn how to visualize networks using the `igraph` package in R.

For this assignment, we will visualize social networking data using anonymized data from Facebook; this data was originally curated in [a recent paper](#) about computing social circles in social networks. In our visualizations, the vertices in our network will represent Facebook users and the edges will represent these users being Facebook friends with each other.

The first file we will use, `edges.csv`, contains variables `V1` and `V2`, which label the endpoints of edges in our network. Each row represents a pair of users in our graph who are Facebook friends. For a pair of friends A and B, `edges.csv` will only contain a single row -- the smaller identifier will be listed first in this row. From this row, we will know that A is friends with B and B is friends with A.

The second file, `users.csv`, contains information about the Facebook users, who are the vertices in our network. This file contains the following variables:

id: A unique identifier for this user; this is the value that appears in the rows of `edges.csv`

gender: An identifier for the gender of a user taking the values A and B. Because the data is anonymized, we don't know which value refers to males and which value refers to females.

school: An identifier for the school the user attended taking the values A and AB (users with AB attended school A as well as another school B). Because the data is anonymized, we don't know the schools represented by A and B.

locale: An identifier for the locale of the user taking the values A and B. Because the data is anonymized, we don't know which value refers to what locale.

PROBLEM 1.1 - SUMMARIZING THE DATA (1/2 points)

Load the data from `edges.csv` into a data frame called `edges`, and load the data from `users.csv` into a data frame called `users`.

EXPLANATION

We can do this operation with:

```
edges = read.csv("edges.csv")
```

```
users = read.csv("users.csv")
```

How many Facebook users are there in our dataset?

✓ Answer: 59

EXPLANATION

From `str(users)` or `nrow(users)`, we see that there are 59 Facebook users in this dataset.

In our dataset, what is the average number of friends per user? Hint: this question is tricky, and it might help to start by thinking about a small example with two users who are friends.

✗ Answer: 4.949153

EXPLANATION

From `str(edges)` or `nrow(edges)`, we see that there are 146 pairs of users in our dataset who are Facebook friends. However, each pair (A, B) must be counted twice, because B is a friend of A and A is a friend of B. To think of this in simpler terms, consider a network with just new people, A and B, and a single edge (A, B). Even though there are two vertices and one edge, each user has on average one friend.

For our network, the average number of friends per user is $292/59=4.95$.

Finally, note that in all likelihood these users have a much higher number of Facebook friends. We are computing here the average number of people in this dataset who are their friends, instead of the average total number of Facebook friends.

You have used 5 of 5 submissions

PROBLEM 1.2 - SUMMARIZING THE DATA (1/1 point)

Out of all the students who listed a school, what was the most common locale?


☐ Locale A☒ Locale B **EXPLANATION**

From `table(users$locale, users$school)`, we read that all students listed at schools A and B listed their locale as B.

You have used 1 of 1 submissions

PROBLEM 1.3 - SUMMARIZING THE DATA (1/1 point)

Is it possible that either school A or B is an all-girls or all-boys school?

☒ No ☐ Yes**EXPLANATION**

We see from `table(users$gender, users$school)` that both genders A and B have attended schools A and B.

You have used 1 of 1 submissions

PROBLEM 2.1 - CREATING A NETWORK (1/1 point)

We will be using the `igraph` package to visualize networks; install and load this package using the `install.packages` and `library` commands.

EXPLANATION

This can be carried out with the following commands:

```
install.packages("igraph")  
  
library(igraph)
```

We can create a new graph object using the `graph.data.frame()` function. Based on `graph.data.frame`, which of the following commands will create a graph `g` describing our social network, with the attributes of each user correctly loaded?

Note: A directed graph is one where the edges only go one way -- they point from one vertex to another. The other option is an undirected graph, which means that the relations between the vertices are symmetric.

☒ `g = graph.data.frame(edges, FALSE, users)` ✓

☐ `g = graph.data.frame(users, FALSE, edges)`

☐ `g = graph.data.frame(edges, TRUE, users)`

☐ `g = graph.data.frame(users, TRUE, edges)`

EXPLANATION

From `graph.data.frame`, we can see that the function expects the first two columns of parameter `d` to specify the edges in the graph -- our `edges` object fits this description.

Our edges are undirected -- if A is a Facebook friend of B then B is a Facebook friend of A. Therefore, we set the directed parameter to `FALSE`.

The vertices parameter expects a data frame where the first column is a vertex id and the remaining columns are properties of vertices in our graph. This is the case with our `users` data frame.

You have used 1 of 1 submissions

PROBLEM 2.2 - CREATING A NETWORK (2/2 points)

Use the correct command from Problem 2.1 to load the graph `g`.

EXPLANATION

Load the graph with:

```
g = graph.data.frame(edges, FALSE, users)
```

Now, we want to plot our graph. By default, the vertices are large and have text labels of a user's identifier. Because this would clutter the output, we will plot with no text labels and smaller vertices:

```
plot(g, vertex.size=5, vertex.label=NA)
```

In this graph, there are a number of groups of nodes where all the nodes in each group are connected but the groups are disjoint from one another, forming "islands" in the graph. Such groups are called "connected components," or "components" for short. How many connected components with at least 2 nodes are there in the graph?

✓ Answer: 4

EXPLANATION

In addition to the large connected component, there is a 4-node component and two 2-node components.

How many users are there with no friends in the network?

✓ Answer: 7

EXPLANATION

There are 7 nodes that are not connected to any other nodes. Each forms a 1-node connected component.

You have used 1 of 3 submissions

PROBLEM 2.3 - CREATING A NETWORK (1/1 point)

In our graph, the "degree" of a node is its number of friends. We have already seen that some nodes in our graph have degree 0 (these are the nodes with no friends), while others have much higher degree. We can use $\text{degree}(g)$ to compute the degree of all the nodes in our graph g .

How many users are friends with 10 or more other Facebook users in this network?

✓ Answer: 9

EXPLANATION

From $\text{table}(\text{degree}(g))$ or $\text{table}(\text{degree}(g) \geq 10)$, we can see that there are 9 users with 10 or more friends in this network.

You have used 2 of 3 submissions

PROBLEM 2.4 - CREATING A NETWORK (2/2 points)

In a network, it's often visually useful to draw attention to "important" nodes in the network. While this might mean different things in different contexts, in a social network we might consider a user with a large number of friends to be an important user. From the previous problem, we know this is the same as saying that nodes with a high degree are important users.

To visually draw attention to these nodes, we will change the size of the vertices so the vertices with high degrees are larger. To do this, we will change the "size" attribute of the vertices of our graph to be an increasing function of their degrees:

$$V(g)\text{size} = \text{degree}(g)/2 + 2$$

Now that we have specified the vertex size of each vertex, we will no longer use the `vertex.size` parameter when we plot our graph:

```
plot(g, vertex.label=NA)
```

What is the largest size we assigned to any node in our graph?

✓ Answer: 11

What is the smallest size we assigned to any node in our graph?

2

✓ Answer: 2

EXPLANATION

From `table(degree(g))` or `summary(degree(g))`, we see that the maximum degree of any node in the graph is 18 and the minimum degree of any node is 0. Therefore, the maximum size of any point is $18/2+2=11$, and the minimum size is $0/2+2=2$.

*You have used 1 of 3 submissions***PROBLEM 3.1 - COLORING VERTICES** (1/1 point)

Thus far, we have changed the "size" attributes of our vertices. However, we can also change the colors of vertices to capture additional information about the Facebook users we are depicting.

When changing the size of nodes, we first obtained the vertices of our graph with `V(g)` and then accessed the size attribute with `V(g)$size`. To change the color, we will update the attribute `V(g)$color`.

To color the vertices based on the gender of the user, we will need access to that variable. When we created our graph `g`, we provided it with the data frame `users`, which had variables `gender`, `school`, and `locale`. These are now stored as attributes `V(g)$gender`, `V(g)$school`, and `V(g)$locale`.

We can update the colors by setting the color to black for all vertices, then setting it to red for the vertices with gender A and setting it to gray for the vertices with gender B:

```
V(g)$color = "black"
```

```
V(g)$color[V(g)$gender == "A"] = "red"
```

```
V(g)$color[V(g)$gender == "B"] = "gray"
```

Plot the resulting graph. What is the gender of the users with the highest degree in the graph?

☐ Missing gender value☐ Gender A☒ Gender B ✓

EXPLANATION

After updating $V(g)$color, run `plot(g, vertex.label=NA)` to plot the graph. All the largest nodes (the ones with the highest degree) are colored gray, which corresponds to Gender B.$

You have used 1 of 1 submissions

PROBLEM 3.2 - COLORING VERTICES (2/2 points)

Now, color the vertices based on the school that each user in our network attended.

Are the two users who attended both schools A and B Facebook friends with each other?

☒ Yes 

☐ No

What best describes the users with highest degree?

☐ None of the high-degree users attended school A

☒ Some, but not all, of the high-degree users attended school A 

☐ All of the high-degree users attended school A

EXPLANATION

As with coloring by gender, we will set the color for all vertices to black, and then we will set the color for students from school A to red and the color for students from schools A and B to gray. Finally we will plot the updated graph:

```
V(g)$color = "black"
```



```
V(g)$color[V(g)$school == "A"] = "red"
```

```
V(g)$color[V(g)$school == "AB"] = "gray"
```

```
plot(g, vertex.label=NA)
```

The two students who attended schools A and B are colored gray; we can see from the graph that they are Facebook friends (aka they are connected by an edge). The high-degree users (depicted by the large nodes) are a mixture of red and black color, meaning some of these users attended school A and other did not.

You have used 1 of 1 submissions

PROBLEM 3.3 - COLORING VERTICES (2/2 points)

Now, color the vertices based on the locale of the user.

The large connected component is most associated with which locale?

☐ Locale A

☒ Locale B ✓

The 4-user connected component is most associated with which locale?

☒ Locale A ✓

☐ Locale B

EXPLANATION

As with the other coloring tasks, we will set the color for all vertices to black, and then we will set the color for users from locale A to red and the color for users from locale B to gray. Finally we will plot the updated graph:

```
V(g)$color = "black"
```

```
V(g)$color[V(g)$locale == "A"] = "red"
```

```
V(g)$color[V(g)$locale == "B"] = "gray"
```

```
plot(g, vertex.label=NA)
```

Nearly all of the vertices from the large connected component are colored gray, indicating users from Locale B. Meanwhile, all the vertices in the 4-user connected component are colored red, indicating users from Locale A.

You have used 1 of 1 submissions

PROBLEM 4 - OTHER PLOTTING OPTIONS (2/2 points)

The help page is a helpful tool when making visualizations. Answer the following questions with the help of `?igraph.plotting` and experimentation in your R console.

Which igraph plotting function would enable us to plot our graph in 3-D?

✓ Answer: rglplot

What parameter to the `plot()` function would we use to change the edge width when plotting `g`?

✓ Answer: edge.width

EXPLANATION

The three functions to plot the igraph are `plot.igraph` (the function we used through the command `"plot"`), `tkplot`, and `rglplot`. `rglplot` makes 3-D plots -- you can try one with `rglplot(g, vertex.label=NA)`. Once you've made the plot, you can click and drag to rotate the graph. To use this function, you will need to install and load the `"rgl"` package.

To change the edge width, you need to change the edge parameter called `"width"`. From `?igraph.plotting`, we read that we need to append the prefix `"edge."` to the beginning for our call to plot, so the full parameter is called `"edge.width"`. For instance, we could plot with edge width 2 with the command `plot(g, edge.width=2, vertex.label=NA)`.

You have used 3 of 4 submissions

Please remember not to ask for or post complete answers to homework questions in this discussion forum.

© All Rights Reserved



© edX Inc. All rights reserved except where noted. EdX, Open edX and the edX and Open EdX logos are registered trademarks or trademarks of edX Inc.

